

Section 4.3 — ARM Time Delay and Instruction Pipeline (Mazidi)

2025-09-01

Chapter 4 • Section 4.3 — Exercises (Mazidi)

Problems are paraphrased to respect copyright. We show the cycle math and the final wall-clock delay using $F = 1/T$.

Timing assumptions used

- One **machine cycle** = one CPU clock period $T = 1/F$.
- Instruction costs used in this section:
NOP = 1, SUBS (reg,#imm) = 1, LDR (literal) = 3, BNE taken = 3, BNE not-taken = 1.
- For a nested delay with outer count **M**, inner count **N**, and **k** NOPs in the inner body, the total cycles are:
TotalCycles = $M \times (k+4) \times N + 5M - 1$
(derivation: inner loop costs = $(k+4)N - 2$; per outer iteration add LDR=3 and the outer SUBS/BNE; include first MOV once).

22) Oscillator frequency if the machine cycle = 1.25 ns

Answer: 800 MHz (i.e., 800 MHz).

23) Machine cycle if $F = 200$ MHz

Answer: 5.00 ns (i.e., 5 ns).

24) Machine cycle if $F = 100$ MHz

Answer: 10.00 ns (i.e., 10 ns).

25) Machine cycle if $F = 160$ MHz

Answer: 6.25 ns (i.e., 6.25 ns).

26) Delay of the subroutine ($M=200$, $N=4,000,000,000$, inner has $k=1$ NOP) at 80 MHz

- **Total cycles:** 4,000,000,000,999
- **Delay:** 50000.000012488 s \approx 13 h 53 min 20.000012 s

27) Delay of the subroutine ($M=100$, $N=50,000,000$, inner has $k=2$ NOPs) at 50 MHz

- **Total cycles:** 30,000,000,499
- **Delay:** 600.000009980 s \approx 10 min 0.000010 s

28) Delay of the subroutine ($M=200$, $N=20,000,000$, inner has $k=3$ NOPs) at 40 MHz

- **Total cycles:** 28,000,000,999
 - **Delay:** 700.000024975 s \approx 11 min 40.000025 s
-

29) Delay of the subroutine ($M=500$, $N=20,000$, inner has $k=3$ NOPs) at 100 MHz

- **Total cycles:** 70,002,499
 - **Delay:** 0.700024990 s \approx 0.700025 s
-

30) “ARM chip does not have the NOP instruction” — what is used instead?

Answer: Assemblers accept the mnemonic **NOP**, which they assemble to a no-effect data-processing instruction such as **MOV r0, r0** (on classic ARM/ARM7). Newer architectures add a real NOP encoding, but on older parts it's this **MOV** pseudo-op.

Cross-checks

- If your core/flash adds wait states, timing will be **longer** than the idealized values above.
- If your toolchain lists different cycle counts (e.g., LDR latency), redo the math with those numbers; the structure stays the same.